

SHREE RAJESH LONDHE

AI/ML ENGINEER

shreelondhe2004@gmail.com | +91 8433593971 | Maharashtra, India

LinkedIn: linkedin.com/in/shree-londhe-b082b2245 | **GitHub:** github.com/Shree-2004 | **Portfolio:** portfolio-bice-nine-ekrh7ppg1i.vercel.app

PROFESSIONAL SUMMARY

AI/ML Engineer shipping Generative AI systems end-to-end — RAG, multi-agent orchestration, and fine-tuned LLMs — then building the guardrails and evals that prove they hold: prompt-injection red-teaming, groundedness scoring, and CI regression gates. Built a hybrid dense+BM25 RAG pipeline scoring 95.6% keyword accuracy and 91.1% groundedness on real financial filings, a 5-agent LangGraph pipeline with citation verification, and two LLM reliability tools — AgentSentinel (5 real defects caught) and AgentLens (pinpoints the exact failing step in a trace). Fine-tuned Mistral-7B via QLoRA, and deployed containerized GenAI microservices as an ML Developer at Qriocity Ventures.

TECHNICAL SKILLS

Programming & Core Tools: Python, C++, C, SQL, Git, GitHub, Pandas, NumPy, Pydantic

Machine & Deep Learning: PyTorch, TensorFlow, Scikit-learn, HF Transformers, PEFT, LoRA, CrossEncoder, bitsandbytes

Generative AI & LLMs: LangGraph, LangChain, LlamaIndex, CrewAI, Multi-Agent Orchestration, OpenAI API, Gemini 2.5, Groq, QLoRA

LLM Evaluation & Guardrails: LLM-as-Judge (calibration, self-consistency, fail-closed design), Prompt-Injection Red-Teaming, Faithfulness & Groundedness Scoring, Hallucination Detection, CI Regression Gating, Ground-Truth Benchmarks, Output-Schema Validation

Storage & Retrieval: Qdrant Cloud, FAISS, PostgreSQL, BM25, Sparse/Dense Hybrid Retrieval

ML Ops & Deployment: FastAPI, Docker, Weights & Biases, Gradio, Streamlit, Railway, Vercel, fpdf6

PROFESSIONAL EXPERIENCE

Qriocity Ventures | ML Developer (Remote)

Jan 2026 – May 2026

- Built end-to-end ML pipelines — preprocessing large-scale datasets, denoising high-frequency signals, and engineering features — improving downstream model training stability and boosting model performance.
- Integrated GenAI frameworks and LLM APIs into client applications using context engineering and vector database indexing, cutting retrieval latency and scaling search across multiple records.
- Deployed containerized production models as Docker/FastAPI microservices, building low-latency API endpoints that cut inference response time and integrated seamlessly with front-end services.

TECHNICAL PROJECTS

SEBI/BSE Financial Filing RAG Analyzer | Qdrant, FastAPI, BM25, PyTorch

github.com/Shree-2004/SEBI-BSE-Financial-Filing-RAG-Analyzer

- Built a hybrid dense + BM25 retrieval pipeline** with CrossEncoder reranking over 3,637 structure-aware chunks from real Reliance/TCS/Infosys filings, scoring 95.6% keyword accuracy, 91.1% groundedness (anti-hallucination) rate, and 97.8% citation coverage on a hand-labeled 45-question eval set.
- Found and fixed a silent chunking bug** causing duplicate content (31,477 → 3,637 chunks) via root-cause debugging; shipped a regression test suite to prevent recurrence.

Multi-Agent Research Assistant | LangGraph, Gemini 2.5 Flash, Tavily, ArXiv

github.com/Shree-2004/Multi-Agent-Research-Assistant

- Architected a 5-agent LangGraph StateGraph pipeline** (Researcher, Analyst, Writer, Verifier, Critic) over shared TypedDict state; a Critic agent scores drafts against a 5-point checklist and routes failures back to the Analyst, capped at one revision cycle by default.
- Built a Citation Verifier** that checks every cited claim against the source text it cites, replacing a heuristic score blind to bad citations; a live 17-source benchmark run measured citation precision falling 1.0 → 0.875 after a Critic-requested revision, while the heuristic score read 10/10 throughout.

AgentSentinel | LangGraph, Google ADK, Streamlit, GitHub Actions

github.com/Shree-2004/Agentsentinel

- Built an LLM guardrails and red-teaming harness** that wraps arbitrary agents behind one adapter interface and scores faithfulness, prompt-injection resistance, tool-call correctness, and latency, gating CI on score regressions.
- Confirmed an injection-resistance verdict on a live agent:** a deliberately poisoned document carrying a hidden instruction-override payload was retrieved as the top match — the agent cited it, answered the real question, and leaked no canary. Faithfulness judge calibrated to 6/6 agreement with hand-labeled cases before any verdict was trusted.
- Audited 3 of my own production agents**, surfacing 5 real defects — including one in the harness itself, where a single scorer's exception discarded every already-collected result; shipped 20 phased commits, 8 pytest suites, and a live Streamlit dashboard.

AgentLens | Python, Gemini, Groq, LLM-as-Judge

github.com/Shree-2004/agentlens

- Built a post-hoc trace debugger** that pairs deterministic rule checks (tool-call loops, ignored errors, missing arguments) with an LLM judge that must extract each step's claims before giving a verdict — treating API failure as JudgeUnavailable, never as “no failure found.”
- Zero false positives across two judge models** on a 10-trace hand-labeled benchmark (80% accuracy on a complete 30-run Groq pass); the fix that caught a real cross-agent contradiction was structural — a required claims array in the judge's output schema — after the prompt-only version missed it 5/5.

Domain Fine-Tuned Finance LLM | Mistral-7B, QLoRA

github.com/Shree-2004/Finance-llm

- Fine-tuned Mistral-7B-Instruct** on custom Indian personal-finance data using QLoRA (4-bit NormalFloat via bitsandbytes), reducing training memory footprint 4× — training loss 1.72 → 0.68 over 3 epochs.

EDUCATION

Dr. Babasaheb Ambedkar Technological University (DBATU) | B.Tech — IT (AI & ML Specialization)

Aug 2026

LANGUAGES: English (Fluent), Hindi (Fluent), Marathi (Native)