Shree Rajesh Londhe

Maharashtra, India | shreelondhe2004@gmail.com | +91 8433593971
github.com/Shree-2004 | linkedin.com/in/shree-londhe | Portfolio

Professional Summary

AI/ML Engineer with production experience deploying containerized GenAI microservices (Docker, FastAPI) and building RAG, multi-agent, and fine-tuned LLM systems. Specializes in LLM evaluation and guardrails — a live financial-filings RAG at 93.5% strict accuracy on 124 PDF-verified questions with cell-level answer verification, and ground-truth benchmarks built to falsify claims rather than confirm them.

Skills

Languages & ML: Python, SQL, C++, PyTorch, TensorFlow, Scikit-learn, HuggingFace Transformers, PEFT, LoRA, QLoRA, Pandas, NumPy
Generative AI: RAG, LangGraph, LangChain, LlamaIndex, CrewAI, Multi-Agent Orchestration, Agentic Workflows, Tool-Calling, Prompt & Context Engineering, OpenAI API, Gemini, Groq
Evaluation & MLOps: LLM-as-Judge, Groundedness Scoring, Hallucination Detection, Ground-Truth Benchmarks, CI Regression Gating, Docker, FastAPI, Weights & Biases, GitHub Actions
Data & Retrieval: Qdrant, FAISS, PostgreSQL, BM25, Hybrid Sparse/Dense Retrieval, CrossEncoder Reranking

Professional Experience

Qriocity Ventures, ML Developer (Remote) *Jan 2026 – May 2026*

- Engineered production ML pipelines — signal denoising, feature engineering, preprocessing — stabilizing client model training runs
- Integrated LLM APIs via context engineering and vector-database indexing, reducing retrieval latency and scaling semantic search
- Deployed containerized models as Docker/FastAPI microservices serving low-latency inference endpoints to front-end clients

Projects

SEBI/BSE Financial Filing RAG Analyzer — Qdrant, FastAPI, Docling, ONNX, Groq / *Live* *GitHub*

- Built cited Q&A over 13 listed companies' SEBI/BSE filings (hybrid dense+BM25, CrossEncoder reranking); found the original 95.6% metric was lenient, re-scored strictly against PDF-verified answers, and traced failures to table parsing — Docling + contextual chunk headers lifted strict accuracy 29% → 87%
- Routed reported-figure, comparison and growth questions to a tool-calling agent (metric lookup, AST-safe calculator): 93.5% strict on 124 questions, 6/6 unanswerable ones refused instead of guessed; cell-level verifier catches 85% of wrong figures (old check: 0%); deployed live in 512 MB via ONNX int8 (~1 GB → 470 MiB)

Stock-judgment-bench — Python, Gemini, yfinance, BSE Filings *GitHub*

- Built a leak-free benchmark testing whether an LLM's stock judgement beats chance across 41 NSE equities and 90 windows (3,690 briefs); 5 leak tests pin point-in-time correctness, one catching a 20× split-adjustment price error
- Scored cross-sectionally so market drift is unexploitable; found no detectable edge (rank corr −0.015, t = −0.65) vs momentum, reversal and random, falsifying a pre-registered hypothesis

Multi-Agent Research Assistant — LangGraph, Gemini, Tavily, ArXiv *GitHub*

- Built a 5-agent LangGraph StateGraph pipeline (Researcher, Analyst, Writer, Verifier, Critic) with a Critic scoring drafts on a 5-point checklist and routing failures back for one capped revision
- Shipped a Citation Verifier that caught precision falling 1.0 → 0.875 on a 17-source run, after a revision the prior heuristic rated 10/10

AgentLens — Python, Gemini, Groq, LLM-as-Judge *GitHub*

- Built an agent trace debugger pairing deterministic rule checks with an LLM judge that must extract each step's claims before ruling; zero false positives across two judge models on a 10-trace benchmark (80% over 30 runs)
- Fixed the judge structurally — a required claims array in its output schema — after the prompt-only version missed a real cross-agent contradiction 5/5

Education

Dr. Babasaheb Ambedkar Technological University (DBATU), Maharashtra *Aug 2026*
B.Tech — Information Technology (AI & ML Specialization) | Languages: English, Hindi, Marathi